

# Lecture 4 Virtual Memory

2025.11.26

# Schedule

- Project 4 assignment
- Project 3 due



# Project 4 Virtual Memory

- Requirement
  - Implement virtual memory management
    - Enable VM for your OS
    - Setup page tables for user-level processes
    - Handle page fault to support on demand paging assuming physical memory is enough
    - Handle page swapping assuming physical memory is limited



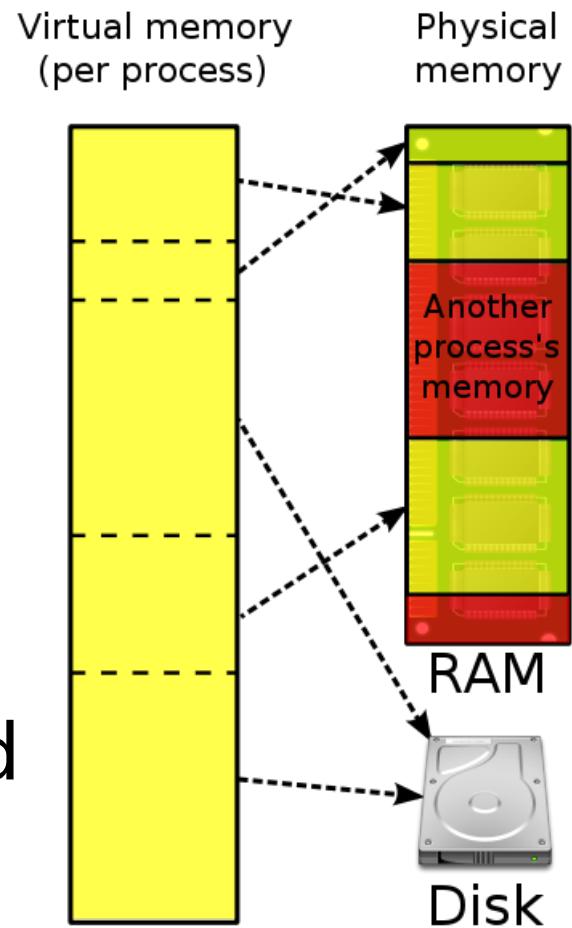
# Project 4 Virtual Memory

- Requirement
  - Implement virtual memory management
    - Support getting free memory
    - Zero-copy based pipe



# Project 4 Virtual Memory

- Virtual memory
  - Each process sees a contiguous and linear address space, which is called virtual addresses
  - Virtual addresses are mapped into physical addresses by both HW and SW



[From Wikipedia]

# Project 4 Virtual Memory

- Virtual memory
  - Virtual address space is divided into pages, which are blocks of contiguous virtual memory addresses
    - e.g. 4KB pages
  - Each page has a virtual address, and is mapped into a physical page frame (e.g. 4KB)

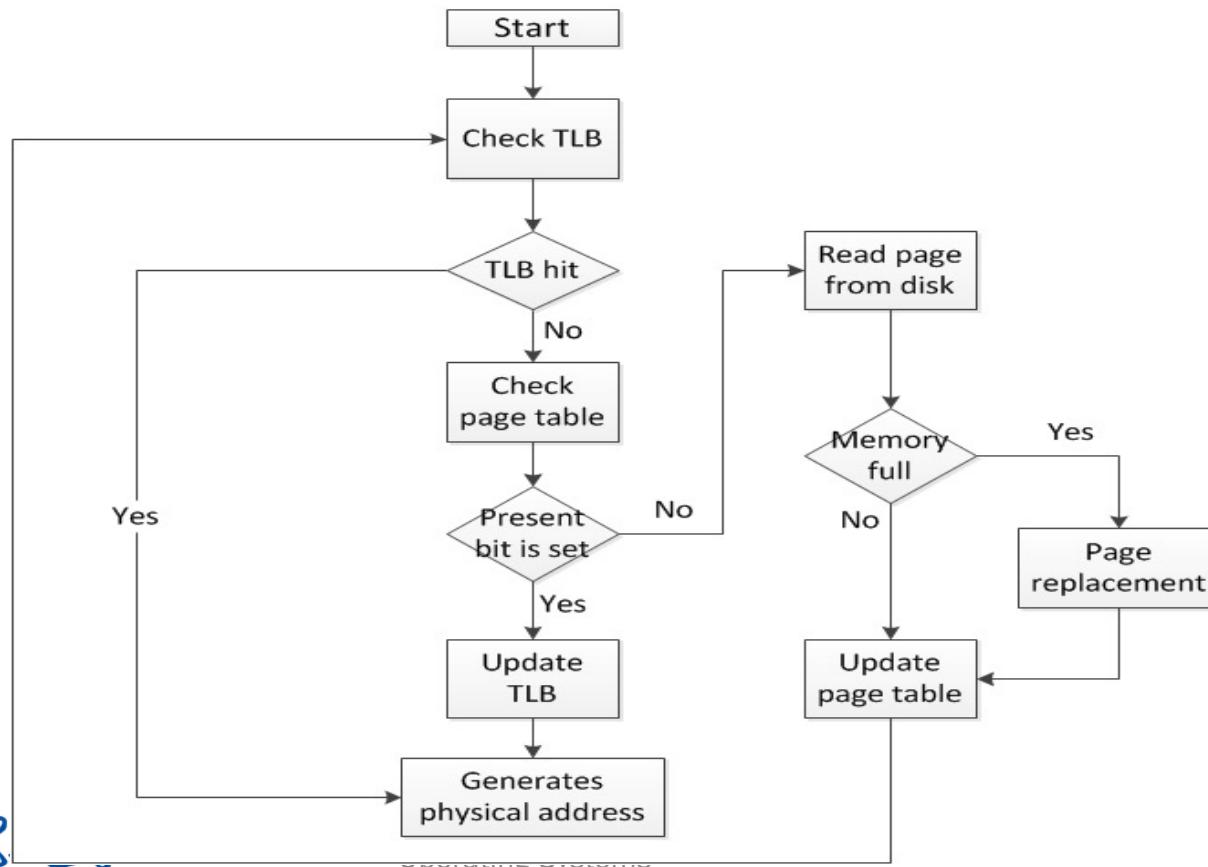


# Project 4 Virtual Memory

- Virtual memory
  - Page tables
    - The data structure to store the mapping between virtual addresses and physical addresses
    - Each mapping is a page table entry
  - MMU and TLB
    - MMU stores a cache of recently used mappings from the page table, which is called translation lookaside buffer (TLB)
  - For RISC-V, TLB is handled by hardware
    - Hardware page table walker

# Project 4 Virtual Memory

- Virtual memory
  - Address translation



# Project 4 Virtual Memory

- Enabling virtual memory
  - You need to set ***satp*** register
    - PPN is the physical page frame of page directory
    - Fill 8 in MODE to enable sv39 mode (three-level page table)
    - ASID represents the ID of current process, which is used to distinguish different processes

63	60 59	44 43	0
MODE (WARL)	ASID (WARL)	PPN (WARL)	
4	16	44	



# Project 4 Virtual Memory

- Enabling virtual memory
  - Example

```
// SATP_MODE_SV39 为 8, ASID 为 0, NORMAL_PAGE_SIZE 为 12, 代表 4KB 的偏移
// 这些宏定义定义在 arch/riscv/include/pgtable.h
set_satp(SATP_MODE_SV39, 0, 0x51000000 >> NORMAL_PAGE_SHIFT);
```



# Project 4 Virtual Memory

- Process switch
  - Modify *satp* register
  - Use *sfence.vma* to flush TLB manually
  - Pls. refer to *local\_flush\_tlb\_all* 和 *local\_flush\_tlb\_page*  
(arch/riscv/include/pgtable.h )
- Note that, when PTE is modified, you also need to flush TLB



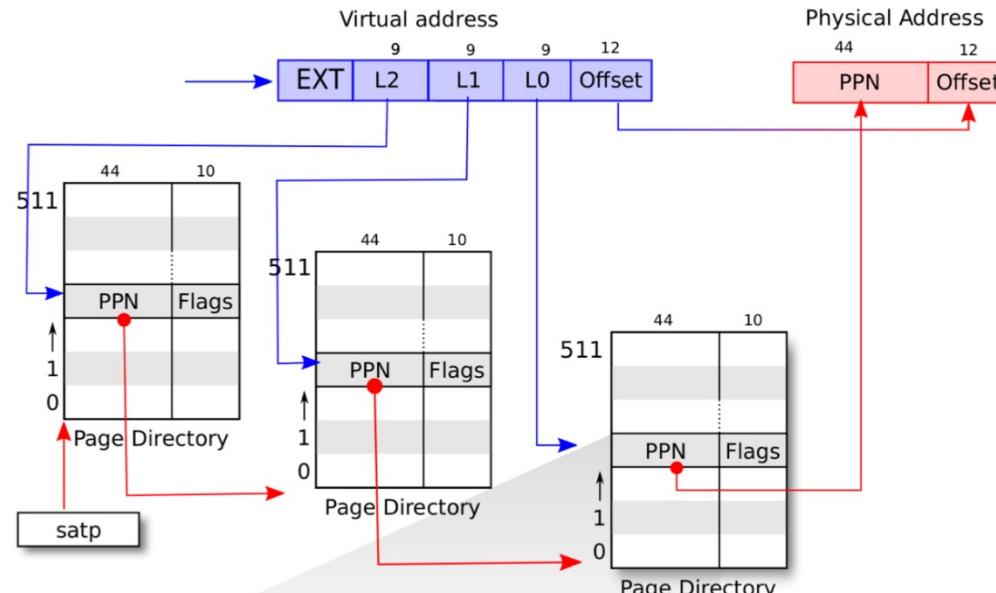
# Project 4 Virtual Memory

- Page table setup
  - After enabling virtual memory, **both kernel and user-level processes** use virtual addresses
  - Both kernel and user-level processes need to setup page tables
  - Pls. note that page table **itself also needs** page frames



# Project 4 Virtual Memory

- Page table setup
  - Two-level page table is used for kernel, which is already setup in start code
  - Three-level page table is used for user-level process



# Project 4 Virtual Memory

- Page table setup
  - Design the structure of page table entries (PTE)
    - Physical address
    - X,W,R
    - U : Set for user-mode access
    - Valid
    - Access, Dirty
      - They are automatically set using QEMU, but need to be set by page fault handler when on board

63	54 53	28 27	19 18	10 9	8	7	6	5	4	3	2	1	0
Reserved	PPN[2]	PPN[1]	PPN[0]	RSW	D	A	G	U	X	W	R	V	

10 26 9 9 2 1 1 1 1 1 1 1 1 1



# Project 4 Virtual Memory

- Page table setup
  - Statically fill page table
    - Fill paired VA to PA mappings
  - On-demand paging
    - An empty page table
    - Fill the page table when page fault occurs



# Project 4 Virtual Memory

- Kernel Address Space
  - The entry point of the kernel is 0xfffffffffc0\_50202000, and its physical address is 0x50202000

---

地址范围	建议用途
0xfffffffffc0_50000000 – 0xfffffffffc0_50200000	BBL 代码及其运行所需的内存
0xfffffffffc0_50200000 – 0xfffffffffc0_51000000	Kernel 的数据段/代码段等
0xfffffffffc0_51000000 – 0xfffffffffc0_52000000	Kernel 页表以及跳转表等
0xfffffffffc0_52000000 – 0xfffffffffc0_60000000	供内核动态分配使用的内核虚拟地址空间

---

# Project 4 Virtual Memory

- Process Address Space
  - Each Process has its private address space
    - Apps may start at same entry point and stack address
  - All processes share kernel address space
    - You need to copy the kernel page table into the page tables of all user-level processes



# Project 4 Virtual Memory

- Load user program
  - First set up page table for user process
  - The kernel needs to load the program to the virtual address
  - The kernel needs to know the **kernel virtual address** mapping to the **corresponding physical address** for the above VA

```
// va 为需要映射的虚拟地址, pgdir 为页表目录,  
// 返回值为为 va 映射的物理地址对应的内核虚地址  
uintptr_t alloc_page_helper(uintptr_t va, uintptr_t pgdir);
```



# Project 4 Virtual Memory

- Load user program

```
1 Elf file type is EXEC (Executable file)
2 Entry point 0x10000
3 There are 2 program headers, starting at offset 64
4
5 Program Headers:
6      Type          Offset        VirtAddr       PhysAddr
7                  FileSiz        MemSiz         Flags  Align
8      LOAD          0x0000000000001000 0x0000000000001000 0x0000000000001000
9                  0x000000000000c35 0x0000000000001c40  RWE      0x1000
10     GNU_STACK    0x0000000000000000 0x0000000000000000 0x0000000000000000
11                  0x0000000000000000 0x0000000000000000  RW      0x10
12
13 Section to Segment mapping:
14   Segment Sections...
15     00  .text .rodata .sdata .sbss .bss
16     01
```

Virtual address range: 0x10000 ~ 0x11c40

Process image location: 0x10000 ~ 0x10c35

Fill zeros: 0x10c35 ~ 0x11c40



# Project 4 Virtual Memory

- Load user program

```
1 Elf file type is EXEC (Executable file)
2 Entry point 0x10000
3 There are 2 program headers, starting at offset 64
4
5 Program Headers:
6      Type          Offset        VirtAddr       PhysAddr
7                  FileSiz        MemSiz         Flags  Align
8      LOAD          0x0000000000001000 0x0000000000001000 0x0000000000001000
9                  0x000000000000c35 0x0000000000001c40  RWE      0x1000
10     GNU_STACK     0x0000000000000000 0x0000000000000000 0x0000000000000000
11                  0x0000000000000000 0x0000000000000000  RW      0x10
12
13 Section to Segment mapping:
14   Segment Sections...
15     00  .text .rodata .sdata .sbss .bss
16     01
```

Virtual address range: 0x10000 ~ 0x11c40

Assume mapped physical address range: 0x52000000 ~ 0x52000c35

Kernel load program image: 0xfffffc0 52000000 ~ 0xfffffc0 52000c35

# Project 4 Virtual Memory

- Handling page fault
  - Page fault: not matching PTE in the page table
    - Allocate a physical page
    - Update the page table to setup the mapping between virtual page and physical page
    - Flush TLB entry
  - Use *scause* register to identify interrupt type
  - Develop page fault handler
  - The address of virtual page causing page fault is stored in *stval* register

# Project 4 Virtual Memory

- Handling swapping
  - Assuming you have limit memory space, you need to swap in-memory pages out to disk when there are no available pages.
  - When these swapped-out pages are reused again, swap the pages in from disk
  - Use bios read/write functions we provided to read from and write to SD card



# Project 4 Virtual Memory

- Step-by-step Task 1
  - Load kernel into memory
  - Finish APIs for manipulating kernel page table
  - Print “CPU #i has entered kernel with VM!”

```
1  /*
2  * Just start kernel with VM and print this string
3  * in the first part of task 1 of project 4.
4  * NOTE: if you use SMP, then every CPU core should call
5  * `kernel_brake()` to stop executing!
6  */
7 printk("> [INIT] CPU #%u has entered kernel with VM!\n",
8     (unsigned int)get_current_cpu_id());
9 // TODO: [p4-task1 cont.] remove the brake and continue to start user
10    processes.
11 kernel_brake();
```



# Project 4 Virtual Memory

- Step-by-step Task 1
  - Statically setup page table for user-processes *shell* and *fly*
    - Initialize page table entries and set all pages valid
    - Filling kernel page table entries into user-level process page table
    - Modify scheduler to allow switch page table among different processes
    - Reclaim memory space when executing *kill* and *exit*

# Project 4 Virtual Memory

- Step-by-step: Task 2
  - Setup dynamic page table for user-level processes to support on-demand paging
  - Handling page fault for *rw*.
    - *rw* runs when you can input different virtual addresses
  - You need support in-kernel locking as you did in Project 2
    - The lock resides in the shared kernel space

# Project 4 Virtual Memory

- Step-by-step: Task 3
  - Limit the available memory space, and design a test case to access a range of memory addresses, which is larger than the available memory space
  - Supporting swapping with your own designed replacement policy
    - We do not have any specific requirement for the replacement policy



# Project 4 Virtual Memory

- Step-by-step: Task 4
  - Support user to use the shell command *free* to check the available memory space
  - Implement the corresponding syscall function `sys_get_free_memory`



# Project 4 Virtual Memory

- Step-by-step: Task 5
  - Support zero copy based pipe
  - After sender sends data into the pipe, the kernel unmaps the mapping between the physical page frames and the virtual pages of the pipe
  - Then, the kernel maps these page frames holding the senders' data to the address space of the receiver
  - After that, the receiver can directly read the data

# Project 4 Virtual Memory

- Requirements for design review
  - 请描述你对内核页表的理解，例如内核页表结构、PTE内容、页表大小、页表在内存中的存放位置等
  - 请描述用户态进程页表创建设置的过程。并简述内核加载一个用户进程时，如何将其运行在虚拟内存上？
  - 请描述page fault处理的流程，最好可以提供一些流程伪代码

# Project 4 Virtual Memory

- Requirements for design review
  - Page swap时使用的替换策略是什么？
  - 任何想讨论的其他问题



# Project 4 Virtual Memory

- Requirement for S/A/C-Core

Core type	Task requirements
S-Core	Task 1
A-Core	Tasks 1~3
C-Core	Tasks 1~5



# Project 4 Virtual Memory

- P4 schedule
  - Design review: 3<sup>rd</sup> Dec.
  - Due
    - C core due : 10<sup>th</sup> Dec.
    - A core due: 17<sup>th</sup> Dec.
  - Please read the guidebook carefully, and contact us if you have any question

