

Lecture 2 A Simple Kernel

2025.10.15


Schedule

- Project 1 due
- Project 2-1 assignment



Project 2 – A Simple Kernel

- Booting procedure



Project 2 – A Simple Kernel

- Requirement I
 - Write a simple kernel (non-preemptive)
 - Process management
 - Start a set of processes
 - Provide context switches between processes with non-preemptive scheduling
 - Support basic mutex locking to enable BLOCK state of processes



Project 2 – A Simple Kernel

- A set of user processes
 - print1, print2, flyer
 - Please first read the codes of different tasks to understand what they do



Project 2 – A Simple Kernel

- Process Control Block (PCB)
 - A in-memory data structure in OS kernel containing the information to manage a process
 - Please refer to the given definition `pcb_t` (*include/os/sched.h*)



Project 2 – A Simple Kernel

- Process Control Block (PCB)
 - Process ID
 - Process status
 - Kernel stack pointer
 - User stack pointer
 - Register context
 - You may need to add new fields to PCB in subsequent projects

```
/* Process Control Block */  
typedef struct pcb  
{  
    /* register context */  
    // NOTE: this order must be preserved, which is defined in regs.h!!  
    reg_t kernel_sp;  
    reg_t user_sp;  
  
    /* previous, next pointer */  
    list_node_t list;  
  
    /* process id */  
    pid_t pid;  
  
    /* BLOCK | READY | RUNNING */  
    task_status_t status;  
  
    /* cursor position */  
    int cursor_x;  
    int cursor_y;  
  
    /* time(seconds) to wake up sleeping PCB */  
    uint64_t wakeup_time;  
} pcb_t;
```




Project 2 – A Simple Kernel

- Initialize PCB
 - Initialize PCBs for tested processes
 - Assign ID, set status/cursor_x/cursor_y
 - PCBs can be organized as a linked list
 - Allocate memory for kernel stack and user stack
 - Please use allocKernelPage/allocUserPage functions
(*kernel/mm/mm.c*)



Project 2 – A Simple Kernel

- Scheduler (non-preemptive kernel)



Project 2 – A Simple Kernel

- Yield
 - A process itself releases the control of CPU by calling kernel scheduler
 - Kernel scheduler places the current running process to the end of the ready queue, and chooses another process to run
 - Please refer to *kernel_yield()* function

Project 2 – A Simple Kernel

- Context switch
 - Kernel scheduler executes context switch to run new process
 - kernel/sched/sched.c *do_scheduler()*

```
void do_scheduler(void)
{
    // TODO: [p2-task3] Check sleep queue to wake up PCBs

    /* **** */
    /* Do not touch this comment. Reserved for future projects. */
    /* **** */

    // TODO: [p2-task1] Modify the current_running pointer.

    // TODO: [p2-task1] switch_to current_running
}
```



Project 2 – A Simple Kernel

- Context switch
 - Save context
 - CPU registers → kernel stack (in memory)
 - Restore context
 - kernel stack → CPU registers



Project 2 – A Simple Kernel

- Kernel stack initialization

- Kernel stack is used to save and restore the context of a process
- Pls. refer to *switchto_context_t*

```
static void init_pcb_stack(
    ptr_t kernel_stack, ptr_t user_stack, ptr_t entry_point,
    pcb_t *pcb)
{
    /* TODO: [p2-task3] initialization of registers on kernel stack
     * HINT: sp, ra, sepc, sstatus
     * NOTE: To run the task in user mode, you should set corresponding bits
     *       of sstatus(SPP, SPIE, etc.).
     */
    regs_context_t *pt_regs =
        (regs_context_t *) (kernel_stack - sizeof(regs_context_t));

    /* TODO: [p2-task1] set sp to simulate just returning from switch_to
     * NOTE: you should prepare a stack, and push some values to
     *       simulate a callee-saved context.
     */
    switchto_context_t *pt_switchto =
        (switchto_context_t *) ((ptr_t) pt_regs - sizeof(switchto_context_t));
}
```



Project 2 – A Simple Kernel

- *switch_to* function
 - Work like a function call, but after the call the return address is another process
 - In *switch_to*, first save the process context of the previous process, and then restore the process context of the next process
 - Where to place the process context?
 - Kernel stack

Project 2 – A Simple Kernel

- *switch_to* function

```
// NOTE: the address of previous pcb in a0
// NOTE: the address of next pcb in a1
ENTRY(switch_to)
    addi sp, sp, -(SWITCH_TO_SIZE)

    /* TODO: [p2-task1] save all callee save registers on kernel stack,
     * see the definition of `struct switchto_context` in sched.h*/

    /* TODO: [p2-task1] restore all callee save registers from kernel stack,
     * see the definition of `struct switchto_context` in sched.h*/

    addi sp, sp, SWITCH_TO_SIZE
    jr ra          新进程的入口地址
ENDPROC(switch_to)
```



Project 2 – A Simple Kernel

- Start a process
 - After initialization, the kernel calls *do_scheduler()*

```
while (1)
{
    // If you do non-preemptive scheduling, it's used to surrender control
    do_scheduler();

    // If you do preemptive scheduling, they're used to enable CSR_SIE and wfi
    // enable_preempt();
    // asm volatile("wfi");
}
```



Project 2 – A Simple Kernel

- Start a process
 - You need to choose a new process to run in *do_scheduler*
 - FIFO is a preferred scheduling algorithm
 - In *do_scheduler()*, *switch_to* is called to perform context switch between the current running process and the new one




Project 2 – A Simple Kernel

- Start a process
 - Note that, when you start the first process, there is no current running user process, how do you handle this?
 - Please refer to pid0_pcb (*include/os/sched.h*)
 - It can be used for store context of kernel process (*main.c*)



Project 2 – A Simple Kernel

- BLOCK state of a process



Project 2 – A Simple Kernel

- Mutex lock
 - What if no process currently holds the lock?
 - Acquire the lock
 - What if the lock is currently held?
 - Wait
 - Implement lock-related functions
 - Manage processes that do not acquire the lock
 - Ready queue vs. wait queue?

Project 2 A Simple Kernel

- Step by step – Task 1
 - Initialize PCBs for tested processes
 - Implement do_scheduler and switch_to
 - Start tested processes and support context switch among these processes as a non-preemptive kernel



Project 2 A Simple Kernel

- Step by step – Task 2
 - Implement mutex lock to support BLOCK state



Project 2 A Simple Kernel

- Requirement for design review
 - *switch_to*函数的工作机制是怎样的？是如何支持两个进程的上下文切换的？
 - 你在初始化PCB和内核栈时需要做哪些事情？
 - 你在switchto_context_t中会保存哪些内容？
 - 请介绍你设计的内核调度器(do_scheduler)的工作流程



Project 2 A Simple Kernel

- Requirement for design review
 - 当一个进程被阻塞时，内核会进行如何的处理？
 - 当一个进程被阻塞或获得资源时，内核会将进程的PCB放置在哪里？
 - 你设计的互斥锁机制能否支持一个进程请求多把锁？



Project 2 – A Simple Kernel

- Requirement for S-Core

Core type	Task requirements
S-Core	Tasks 1, 2



Project 2 – A Simple Kernel

- P2 schedule
 - 22nd Oct.
 - P2-1 review
 - P2-2 assign
 - 29th Oct.
 - P2-1 due

