

Lecture 2 A Simple Kernel

2025.10.29

Schedule

- Project 2 part2 assignment
- Project 2 part1 due



Project 2 – A Simple Kernel

- Requirement II
 - Write a simple kernel (preemptive)
 - Provide preemptive scheduling supporting exception handler, including
 - System call handler
 - Clock interrupt handler
 - Round-robin scheduler
 - A complex scheduler (optional for C-Core)



Project 2 A Simple Kernel

- A set of user process
 - User processes call syscalls to invoke kernel services
 - Please refer to unistd.h(tiny_libc/include) and syscall.c (tiny_libc/)

```
void sys_sleep(uint32_t time);  
void sys_yield(void);  
void sys_write(char *buff);  
void sys_move_cursor(int x, int y);  
void sys_reflush(void);  
long sys_get_timebase(void);  
long sys_get_tick(void);  
int sys_mutex_init(int key);  
void sys_mutex_acquire(int mutex_idx);  
void sys_mutex_release(int mutex_idx);
```



Project 2 A Simple Kernel

- Interrupt
 - A signal to the processor emitted by hardware indicating an event requiring immediate process
 - The processor responds by suspending its current running task, saving its state, and executing interrupt handler
 - After the interrupt handler finishes, the processor resumes normal activities


Project 2 A Simple Kernel

- Interrupt

Normal execution



**Normal execution
with interrupt**



中国科学院


University of Chinese Academy of Sciences

Project 2 A Simple Kernel

- Hardware interrupt
 - A change in execution caused by a hardware event
 - Timer interrupt for timesharing
 - I/O device interrupt: disk, network, keyboard etc.
- We do not handle other interrupts here
 - Page fault, TLB miss etc.
- Trap
 - System calls

Project 2 A Simple Kernel

- Handling exception
 - Save context
 - Determine what causes exception
 - Invoke specific routine based on type of exception
 - Restore context
 - Return from exception



Project 2 A Simple Kernel

- Control Status Registers(CSRs)
 - We need to manipulate CSRs when implementing interrupt handler
- Two types of CSRs
 - One type is for setting up interrupt
 - The other is for handling interrupt
 - Pls. refer to Table P2-2 in guidebook



Project 2 A Simple Kernel

- Control Status Registers(CSRs)

Number	Privilege	Name	Description
Supervisor Trap Setup			
0x100	SRW	sstatus	Supervisor status register.
0x102	SRW	sedeleg	Supervisor exception delegation register.
0x103	SRW	sideleg	Supervisor interrupt delegation register.
0x104	SRW	sie	Supervisor interrupt-enable register.
0x105	SRW	stvec	Supervisor trap handler base address.
0x106	SRW	scounteren	Supervisor counter enable.
Supervisor Trap Handling			
0x140	SRW	sscratch	Scratch register for supervisor trap handlers.
0x141	SRW	sepc	Supervisor exception program counter.
0x142	SRW	scause	Supervisor trap cause.
0x143	SRW	stval	Supervisor bad address or instruction.
0x144	SRW	sip	Supervisor interrupt pending.
Supervisor Protection and Translation			
0x180	SRW	satp	Supervisor address translation and protection.



Project 2 A Simple Kernel

- Setup interrupt
 - sie register
 - STIE(Supervisor-level timer interrupt enable)
 - Use enable_preempt
(arch/riscv/kernel/entry.S) to set up timer interrupt

SXLEN-1	10	9	8	6	5	4	2	1	0
WPRI		SEIE	WPRI	STIE	WPRI	SSIE	WPRI		
SXLEN-10	1	3	1	3	1	1	1	1	



Project 2 A Simple Kernel

- Setup interrupt
 - sstatus register

SXLEN-1	SXLEN-2		34	33	32	31		20	19	18	17
SD		WPRI		UXL[1:0]		WPRI		MXR	SUM	WPRI	
1		SXLEN-35		2		12		1	1	1	1
16	15	14	13	12	9	8	7	6	5	4	0
XS[1:0]	FS[1:0]	WPRI	SPP	WPRI	UBE	SPIE	WPRI	SIE	WPRI		
2	2	4	1	1	1	1	3	1	1		



Project 2 A Simple Kernel

- Setup interrupt
 - SIE (Supervisor Interrupt Enable) bit enables or disables all interrupts in supervisor mode
 - SPIE (Supervisor Previous Interrupt Enable) bit indicates whether supervisor interrupts were enabled prior to trapping into supervisor mode.
 - SPP (Supervisor Previous Privilege Mode) bit indicates the privilege mode before the trapping
 - You need to correctly initialize SPIE(1), SPP(0) in sstatus register



Project 2 A Simple Kernel

- Manipulate CSR registers
 - csrr: read csr
 - Loads data from a CSR register into a CPU register
 - csrw: write csr
 - Stores data into a CSR register
 - csrc/csrs: clear/set a CSR register' s corresponding bits
- Examples
 - csrr a0, sip
 - csrw stvec, t0



Project 2 A Simple Kernel

- What if interrupt occurs while in interrupt handler?
 - The processor automatically disables all interrupts when an interrupt occurs
 - The processor automatically re-enables all interrupts after sret is called



Project 2 A Simple Kernel

- Entry to handle exception
 - stvec register
 - BASE field hold the exception handler' s address
 - exception_handler_entry in entry.S is the entry point for handling exceptions



Value	Name	Description
0	Direct	All exceptions set pc to BASE.
1	Vectored	Asynchronous interrupts set pc to BASE+4×cause.
≥2	—	<i>Reserved</i>



Project 2 A Simple Kernel

- Identifying exception type
 - scause register
 - pls. refer to Table P2-3
 - When exception occurs, uses scause register to identify exception type

XLEN-1	XLEN-2	0
Interrupt	Exception Code (WLRL)	
1	XLEN-1	

Interrupt	Exception Code	Description
1	0	<i>Reserved</i>
1	1	Supervisor software interrupt
1	2–4	<i>Reserved</i>
1	5	Supervisor timer interrupt
1	6–8	<i>Reserved</i>
1	9	Supervisor external interrupt
1	10–15	<i>Reserved</i>
1	≥16	<i>Available for platform use</i>
0	0	Instruction address misaligned
0	1	Instruction access fault
0	2	Illegal instruction
0	3	Breakpoint
0	4	Load address misaligned
0	5	Load access fault
0	6	Store/AMO address misaligned
0	7	Store/AMO access fault
0	8	Environment call from U-mode
0	9	Environment call from S-mode
0	10–11	<i>Reserved</i>



Project 2 A Simple Kernel

- Return from exception
 - sepc register
 - The PC address when exception occurs is stored in sepc
 - sret can return to the address in sepc
 - Please pay attention to the sepc address when handling syscall
 - sret should return to sepc+4



Project 2 A Simple Kernel

- Support process sleep
 - Blocking sleep
 - Block the task when it calls sleep()
 - Use a separate queue to keep sleeping tasks
 - Wake up the task
 - When the timing reaches sleeping threshold of the task
 - About timing
 - We provide timing related functions in time.c
 - Use these functions to decide whether sleeping ends

Project 2 A Simple Kernel

- Implement timer interrupt handler
 - Deal with normal tasks
 - Schedule based on your scheduling policy
 - Round Robin policy is OK
 - Deal with sleeping tasks
 - Check whether waking up the task
 - Reset timer using bios_set_timer
(include/os/bios.h)
 - CPU frequency is obtained using bios_read_fdt in main.c

Project 2 A Simple Kernel

- Step by step – Task3
 - Implement syscalls
 - Support do_sleep
 - In this task, we do not have timer interrupt, you need to consider when to check sleeping time



Project 2 A Simple Kernel

- Step by step – Task 4
 - Implement timer interrupt handler
 - Implement round robin scheduler



Project 2 A Simple Kernel

- Step by step – Task 5
 - There exist five airplanes with different moving speed
 - Please design a scheduling policy to allow the five airplanes reach the checkpoint simultaneously as well as reach the end simultaneously
 - e.g. the slow airplane gets longer time slice than the fast one

-=\\"\\
-=\\"c`~~~~~) /~~`
`~~~~~ / /
-==/ /
'_'



A note to “loadbootd”

- Open CPU in DASIC mode
 - Prevent user code jump/call kernel code
 - Prevent user code access kernel data
 - Provide memory protection before virtual memory is available
 - Run both in QEMU and Pynq
- A/C core should pass loadbootd check



Project 2 A Simple Kernel

- Requirement for design review
 - 请描述syscall的执行流程，以及开发syscall时，如何实现参数传递？
 - 在初始化中断处理时，需要设置哪些寄存器，具体初始化为什么值？
 - 当进程因异常进入内核时，需要进行哪些上下文保存？
 - 开发任务3时，init_pcb_stack函数需要完成哪些初始化操作？
 - 任务3和4中，分别何时唤醒睡眠的进程？
 - Any questions you are not clear enough



Project 2 – A Simple Kernel

- Requirement for S/A/C-Core

Core type	Task requirements
S-Core	Tasks 1, 2,3- (只用打印例外报错信息即可)
A-Core	Tasks 1,2,3,4
C-Core	Tasks 1,2,3,4,5



Project 2 – A Simple Kernel

- P2-2 schedule
 - 5th Nov.
 - P2 part2 design review
 - 12th Nov.
 - P2 part2 due

