

Lecture 5 Device Driver

2025.12.10

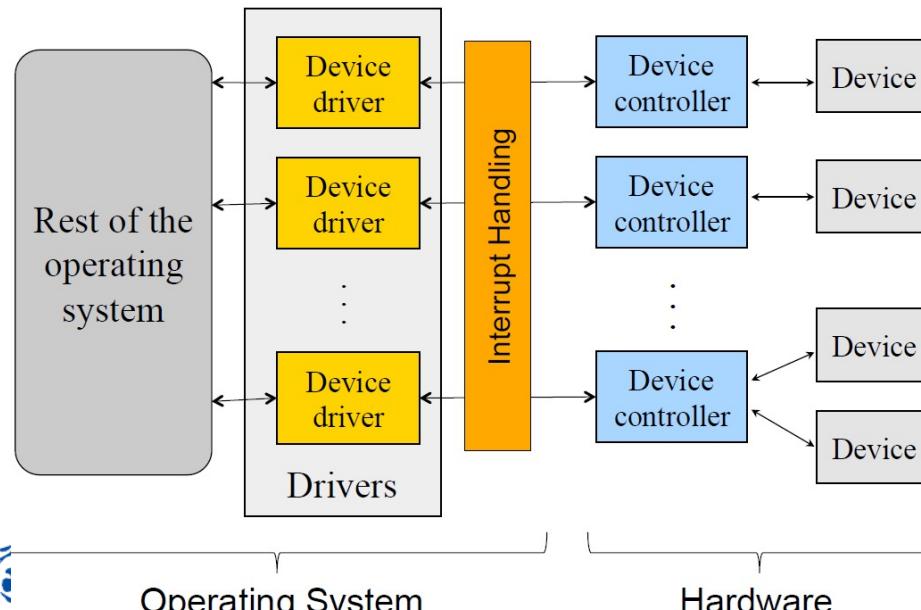
Project 5 Device Driver

- Requirement
 - Implement driver for network card
 - Setup MAC controller registers to allow network card initialization, and sending/receiving data without interrupt
 - Implement network interrupt handler to send/receive data



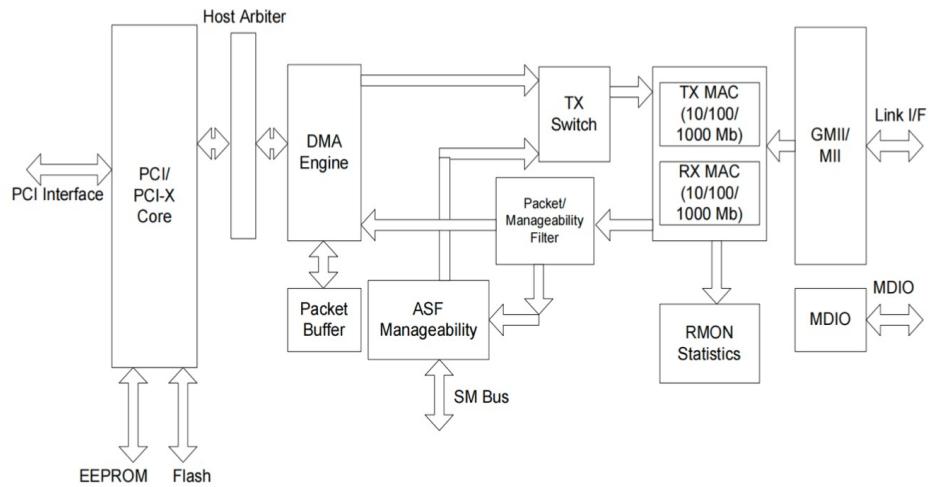
Project 5 Device Driver

- Device driver
 - Interact with device controller
 - Issue commands to device controller to finish data reads and writes
 - e.g. network and disk



Project 5 Device Driver

- MAC controller for E1000 NIC
 - Physical layer
 - Data link layer
 - MAC controller
 - Networking layer
 - Transportation layer
 - Application layer



Project 5 Device Driver

- MAC controller
 - MAC registers
 - register group, including a number of registers
 - *bios_read_fdt()* returns the base address of the register group (physical address)
 - ioremap
 - You need to map the above physical address to kernel virtual address
 - You need to flush TLB after finishing ioremap
 - Accessing a register
 - Base address + $i * 4$
 - i indicates i^{th} register (i starts from 0)



Project 5 Device Driver

- Direct Memory Access
 - Exchange data between host and device bypassing CPU
 - DMA transmit descriptor
 - DMA receive descriptor
 - Use physical addresses



Project 5 Device Driver

- DMA descriptor
 - Each descriptor is 16 bytes
 - Transmit descriptor for transmitting data
 - Receive descriptor for receiving data
 - One descriptor for one network frame



Project 5 Device Driver

- DMA descriptor - Tx descriptor
 - Buffer address: **physical address** of the buffer for transmitting data
 - Length: data length

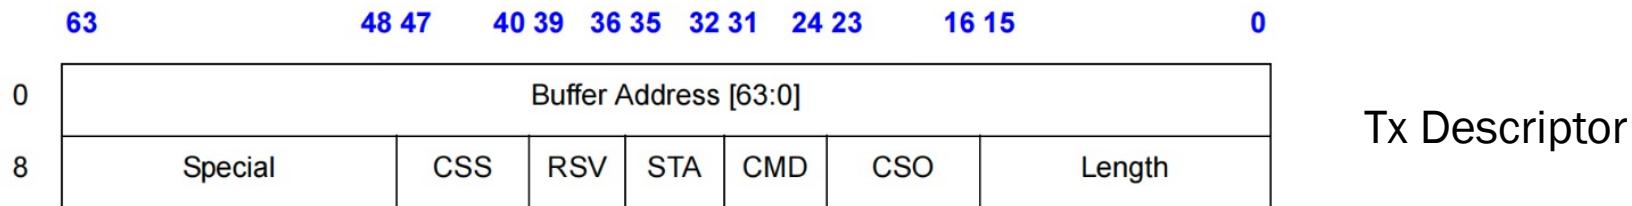
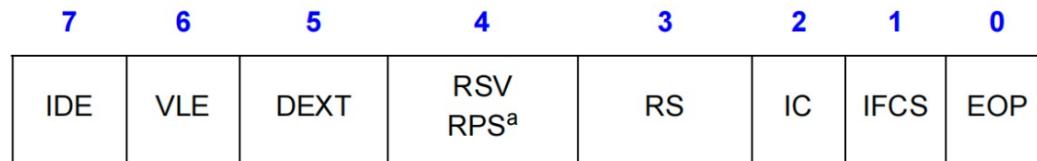


图 P5-3: E1000 发送描述符 (Legacy)

Project 5 Device Driver

- DMA descriptor - Tx descriptor
 - CMD
 - DEXT: 0, indicates Legacy format
 - RS: 1, DMA controller records transmission status in STA field
 - EOP: indicates whether the current descriptor is the last one of the current packet



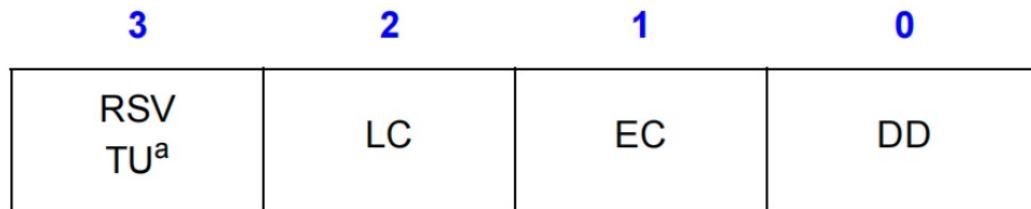
a. 82544GC/EI only.

图 P5-4: E1000 发送描述符 CMD 字段



Project 5 Device Driver

- DMA descriptor - Tx descriptor
 - STA
 - DD: indicates the transmission of current frame is finished



a. 82544GC/EI only.

图 P5-5: E1000 发送描述符 STA 字段



Project 5 Device Driver

- Circular buffer for Tx descriptors
 - TDBAL, TDBAH registers: the low and high 32-bits of the address of the circular buffer
 - TDLEN register: the bytes of the circular buffer
 - Head: the position where DMA controller is currently processing
 - Tail: the position where OS is currently processing

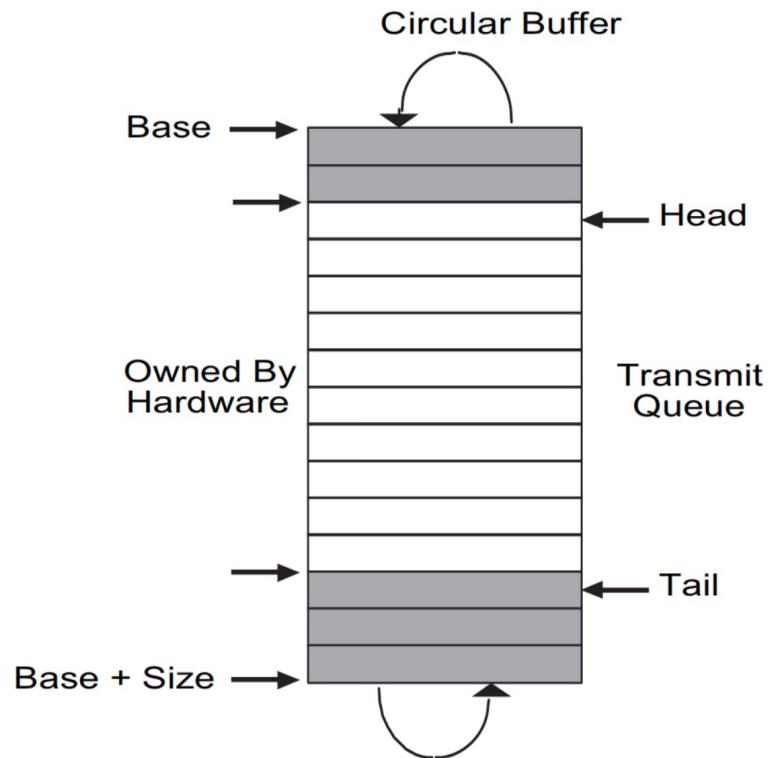


图 P5-6: E1000 发送描述符循环数组

Project 5 Device Driver

- Circular buffer for Tx descriptors
 - TDH register: stores the index of the descriptor element pointed by Head in the circular buffer
 - TDT register: stores the index of the descriptor element pointed by Tail in the circular buffer



Project 5 Device Driver

- Enabling data transmission
 - TCTL register: 1, enable data transmission
 - Once TDH is not equals TDT, DMA controller starts to transmit data



Project 5 Device Driver

- DMA descriptor - Rx descriptor
 - Buffer address: **physical address** of the buffer for receiving data
 - Length: data length

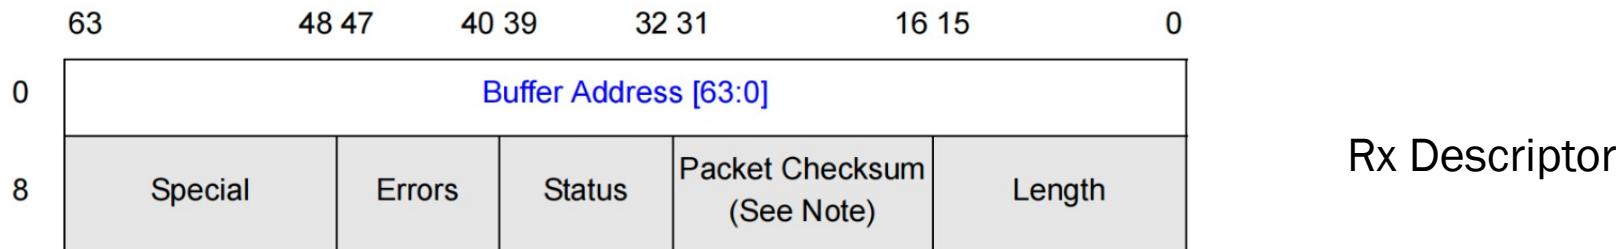


图 P5-13: E1000 接收描述符

Project 5 Device Driver

- DMA descriptor - Rx descriptor
 - STA (Status)
 - EOP: indicates whether the current descriptor is the last one
 - DD: indicates current frame is finished

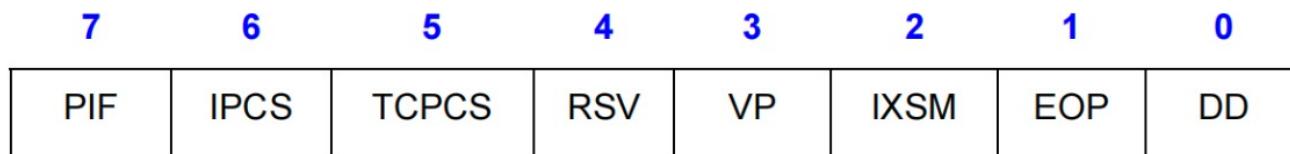


图 P5-14: E1000 接收描述符 STA 字段

Project 5 Device Driver

- Circular buffer for Rx descriptors
 - RDBAL, RDBAH registers: the low and high 32-bits of the circular buffer
 - RDLEN register: the bytes of the circular buffer
 - Head: the position where DMA controller is currently processing
 - Tail: the position where OS is currently processing

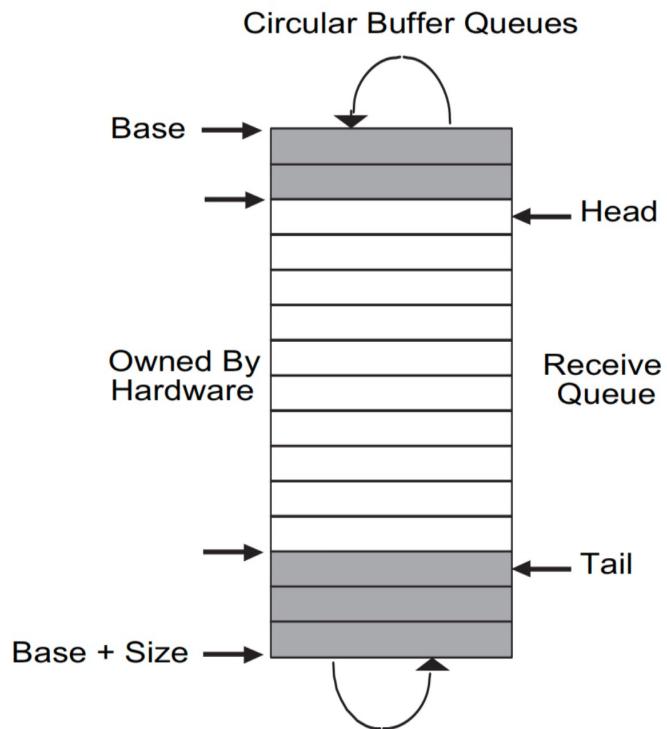


图 P5-15: E1000 接收描述符循环队列

Project 5 Device Driver

- Circular buffer for Rx descriptors
 - RDH register: stores the index of the descriptor element pointed by Head in the circular buffer
 - RDT register: stores the index of the descriptor element pointed by Tail in the circular buffer



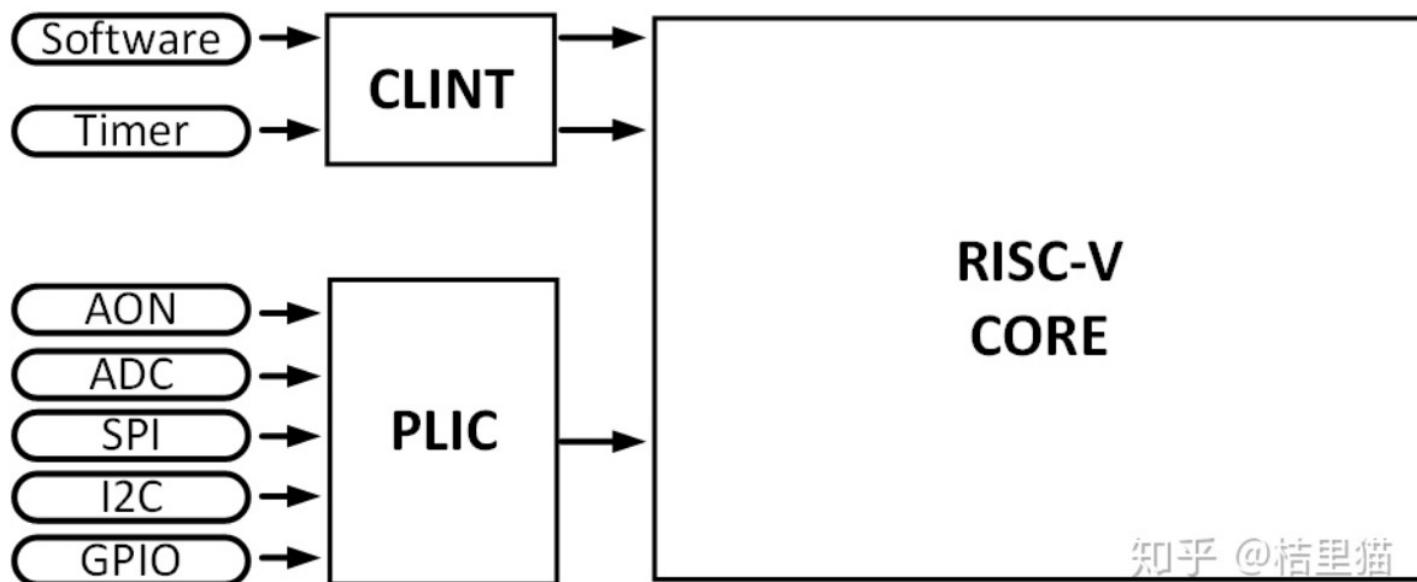
Project 5 Device Driver

- Enabling data receiving
 - RCTL register: 1, enable data receiving
 - NIC only receives the packets with the matched MAC address
 - You need to fill MAC address into RAL0 and RAH0 registers before receiving data



Project 5 Device Driver

- MAC interrupt handler
 - PLIC(Platform Level Interrupt Controller)

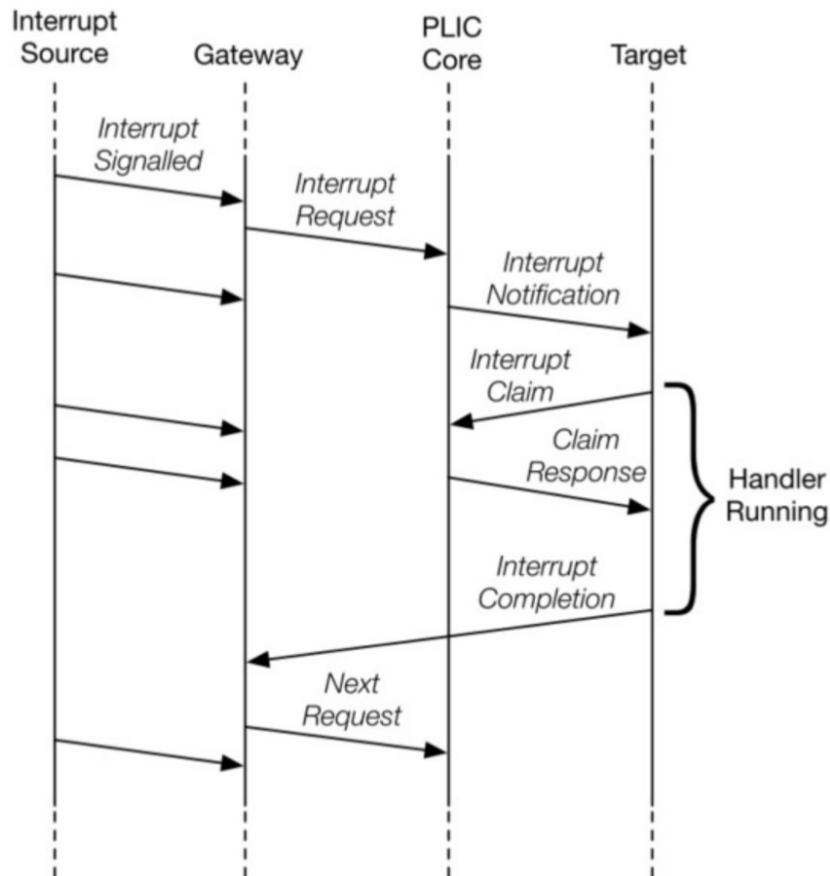


Project 5 Device Driver

- MAC interrupt handler
 - Use scause register to judge IRQ_S_EXT interrupt
 - Use registers provided by PLIC to judge and process NIC interrupt
 - claim/complete register
 - read device ID from claim register and indicate accepting the interrupt from PLIC
 - » NIC device ID: 33 (qemu), 3(board)
 - write device ID to complete register to indicate finishing interrupt handling

Project 5 Device Driver

- MAC interrupt handler



Project 5 Device Driver

- MAC interrupt handler
 - Use registers provided by PLIC to judge and process NIC interrupt
 - IMS register (Interrupt Mask Set): write 1 to it to enable NIC interrupt
 - IMC register (Interrupt Mask Clear): write 1 to it to clear the corresponding bit in IMS register, then NIC interrupt can be disabled



Project 5 Device Driver

- MAC interrupt handler
 - Use registers provided by PLIC to judge and process NIC interrupt
 - ICR register (Interrupt Cause Read): when NIC interrupt occurs, the corresponding bit in ICR is set
 - ICS register (Interrupt Cause Set): write 1 to it then the corresponding bit in ICR is also set
 - We do not use this register in P5



Project 5 Device Driver

- MAC interrupt handler
 - Use registers provided by PLIC to judge and process NIC interrupt
 - TXQE interrupt for data transmission
 - When TDH equals to TDT, this interrupt is triggered
 - RXDMT0 interrupt for receiving data
 - Set RCTL.RXDMT to 00
 - When the Rx descriptors owned by NIC is less than half, this interrupt is triggered

Project 5 Device Driver

- Step-by-step: Task1
 - Initialize DMA descriptors for transmitting data
 - Set corresponding DMA registers to allow transmitting data
 - Implement the syscall for sending data through NIC
 - The number of packets you send needs to be larger than the number of Tx descriptors in the circular buffer

Project 5 Device Driver

- Step-by-step: Task2
 - Initialize DMA descriptors for receiving data
 - Set corresponding DMA registers to allow receiving data
 - Implement the syscall for receiving data through NIC
 - The number of packets you receives needs to be larger than the number of Rx descriptors in the circular buffer

Project 5 Device Driver

- Step-by-step: Task 3
 - Enable MAC interrupt
 - Implement MAC interrupt handler for receiving data and transmitting data
 - The receive thread blocks itself when there is less than half of total Rx descriptors.
 - The transmit thread blocks itself after all descriptors are used. Once NIC finishes transmission, it is waked up



Project 5 Device Driver

- Step-by-step: Task 4
 - Design and implement a new syscall *sys_net_recv_stream* to receive n bytes in sequence
 - You need to handle packet loss or out-of-order packet arriving
 - Please refer to Sec. 10 in the guidebook to implement a reliable transfer protocol, supporting DAT, ACK and RSD packet

Project 5 Device Driver

- Requirement for S/A/C-Core

Core type	Task requirements
S-Core	Task 1, 2
A-Core	Tasks 1~3
C-Core	Tasks 1~4



Project 5 Device Driver

- Task testing
 - For transmitting data
 - Use wireshark (Windows) or tcpdump (Linux) to examine if the transmit thread successfully transmit data
 - Please read our reference documents
 - For receiving data
 - Use pktRxTx to transmit data from host to your development card to test receiving data

Project 5 Device Driver

- Requirements for design review
 - 请分别展示发送和接收DMA描述符的示例，并介绍所填充字段的含义
 - 请介绍网卡中断的处理流程，包含使用哪些寄存器、如何使能、判断和处理中断
 - 任何想讨论的问题



Project 5 Device Driver

- P5 schedule
 - Design review
 - 17th Dec.
 - C core due
 - 24th Dec.
 - A core due
 - 31st Dec.



Project 5 Device Driver

- P6 schedule
 - Assignment
 - 24th Dec.
 - Design review
 - 31st Dec.
 - C core due
 - 7th Jan.
 - Final due
 - 14th Jan.

